FlashRAG

Blazing-fast RAG inference for the impatient

Vitaly Kleban, vk@cyber.fund, @vkleban

Naïve RAG Inference

Naïve RAG Inference

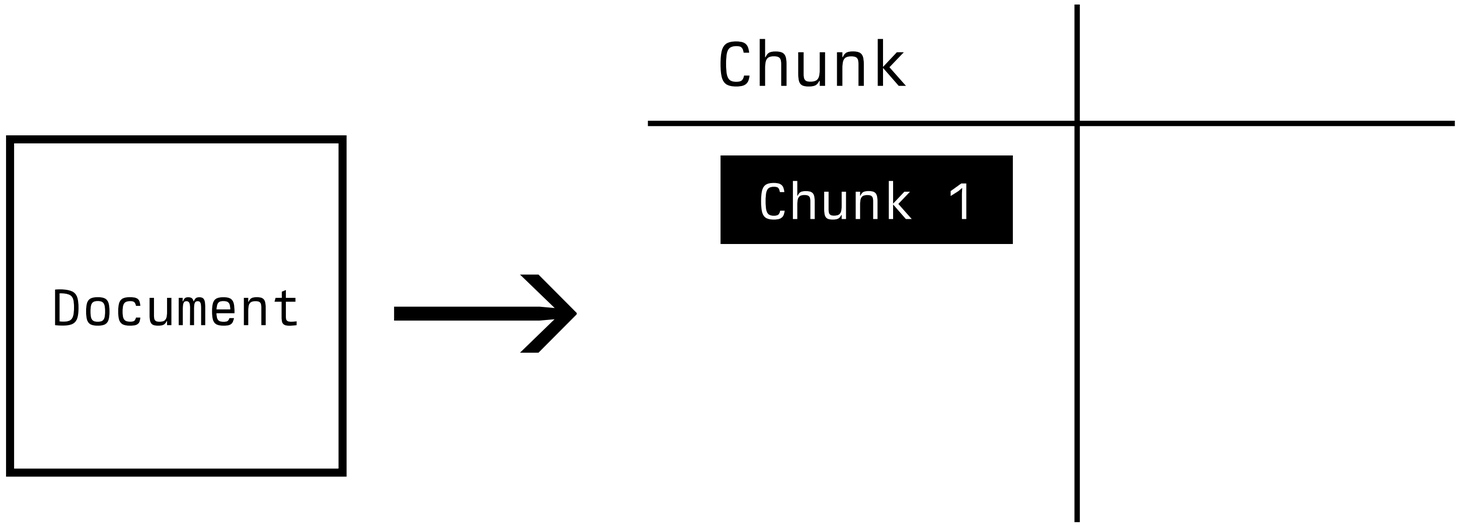

Naïve RAG Inference

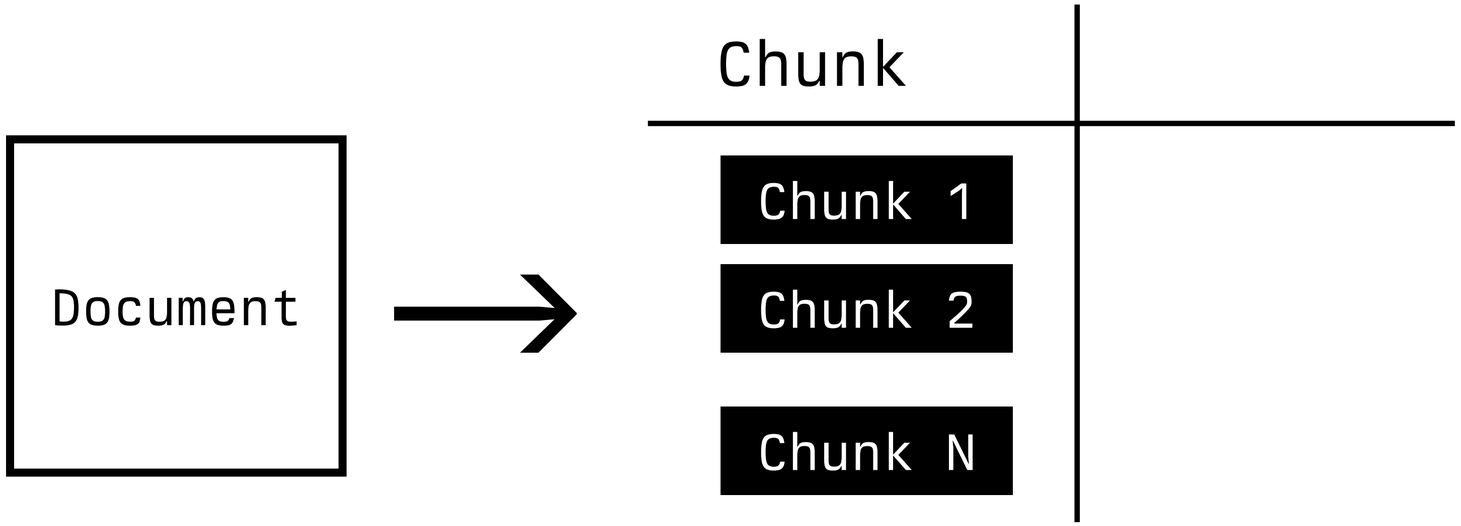

Naïve RAG Inference

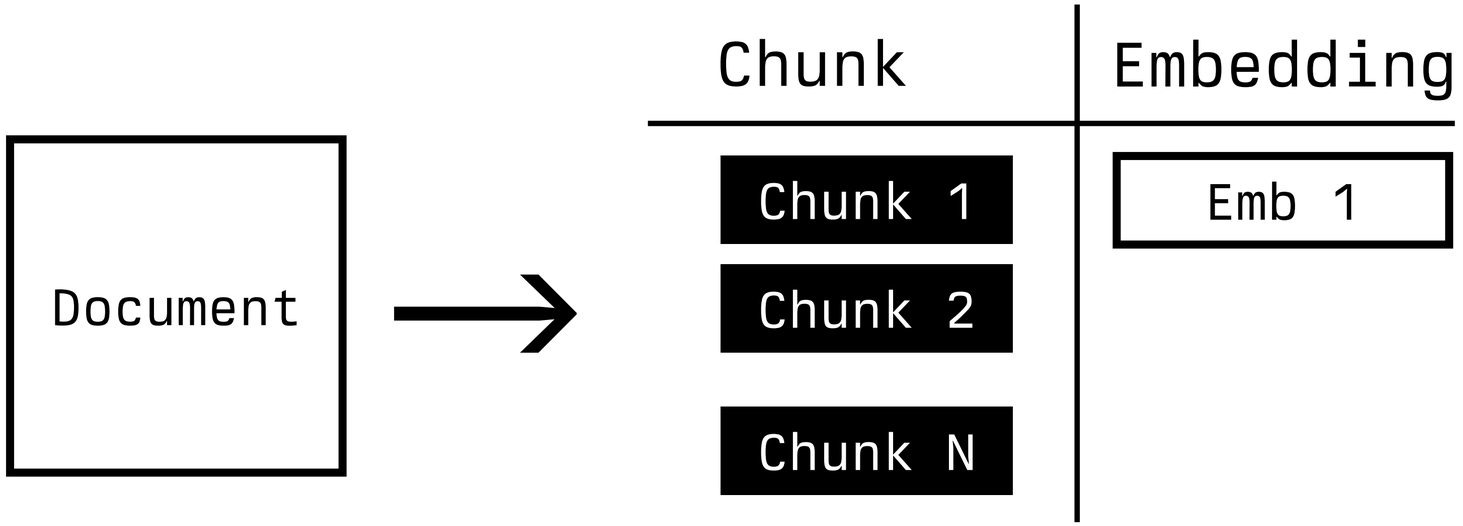

Naïve RAG Inference

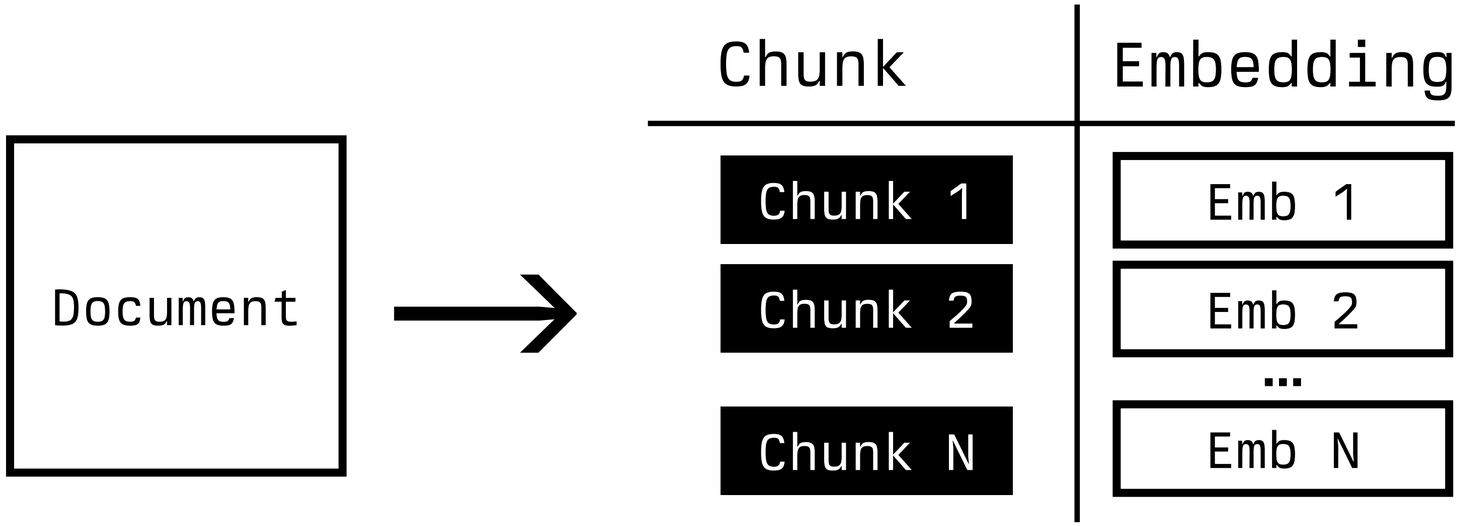

Naïve RAG Inference

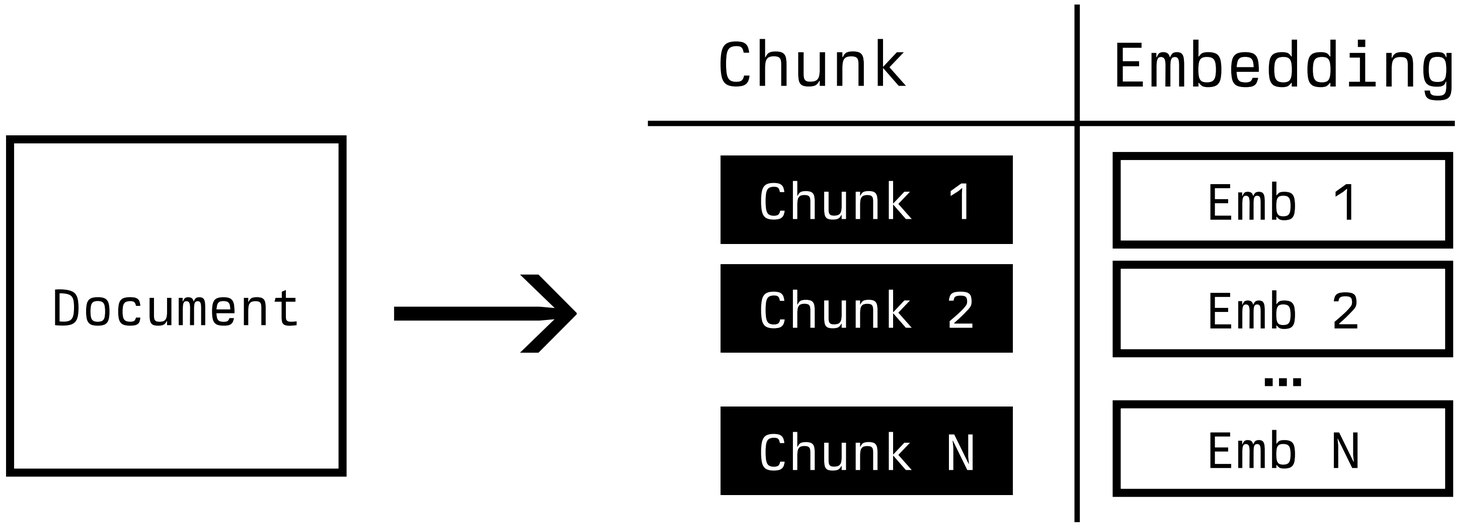

Naïve RAG Inference

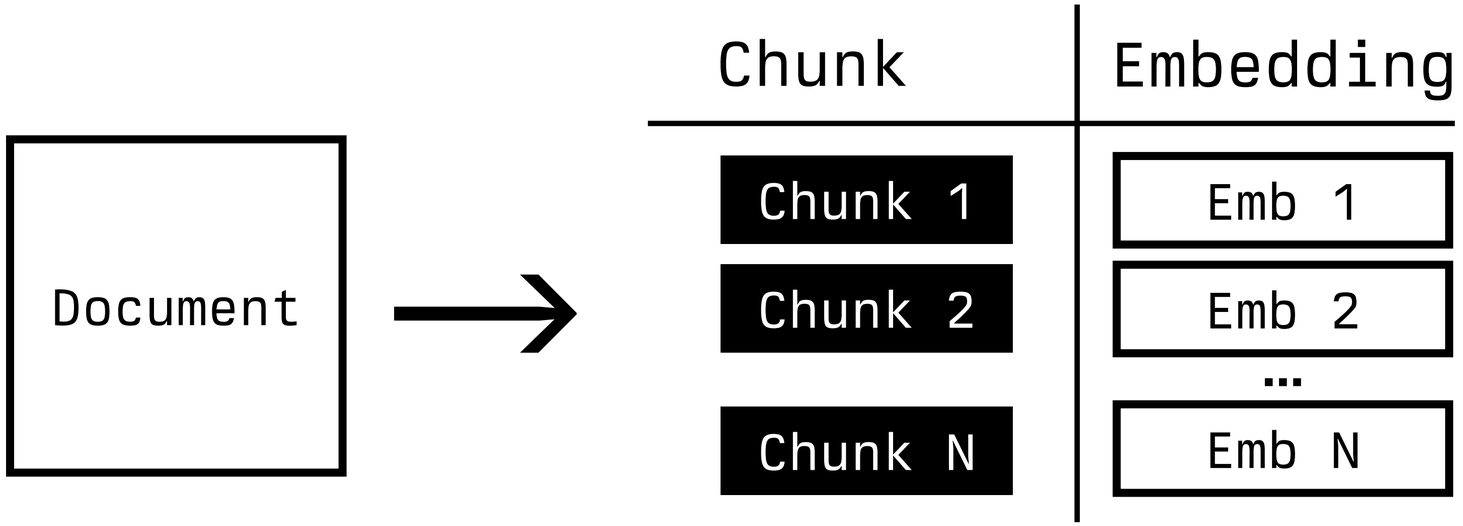

Naïve RAG Inference

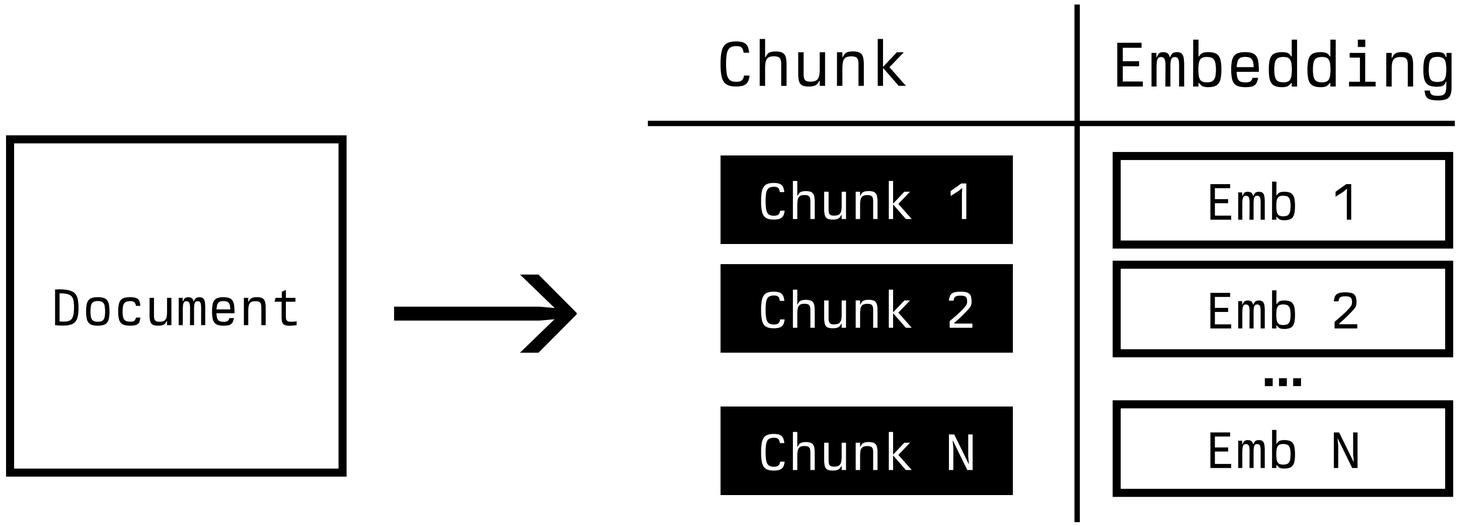

Flash RAG Inference

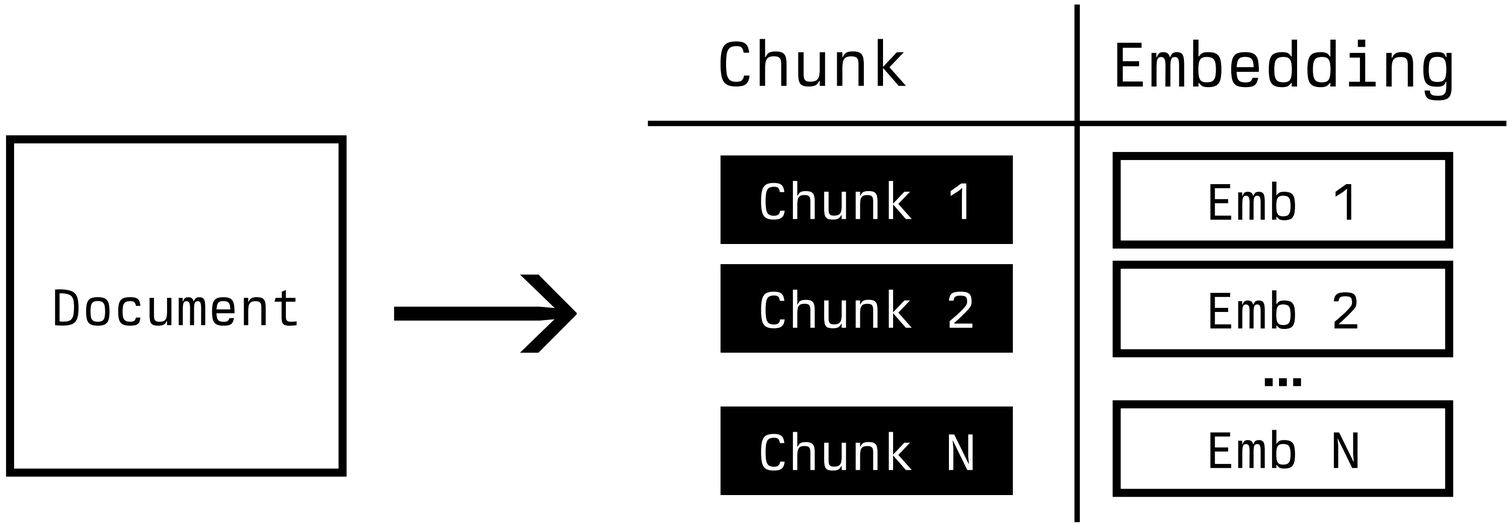

Flash RAG Inference

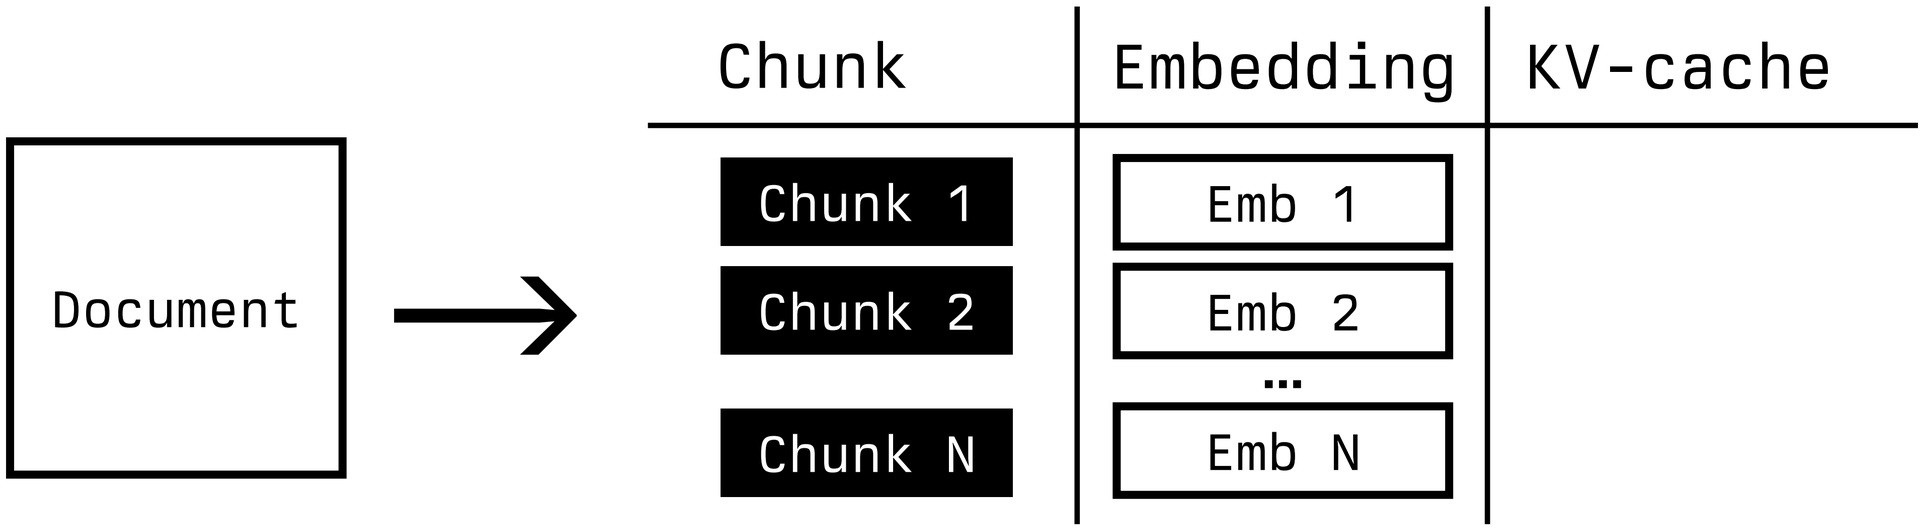

Flash RAG Inference

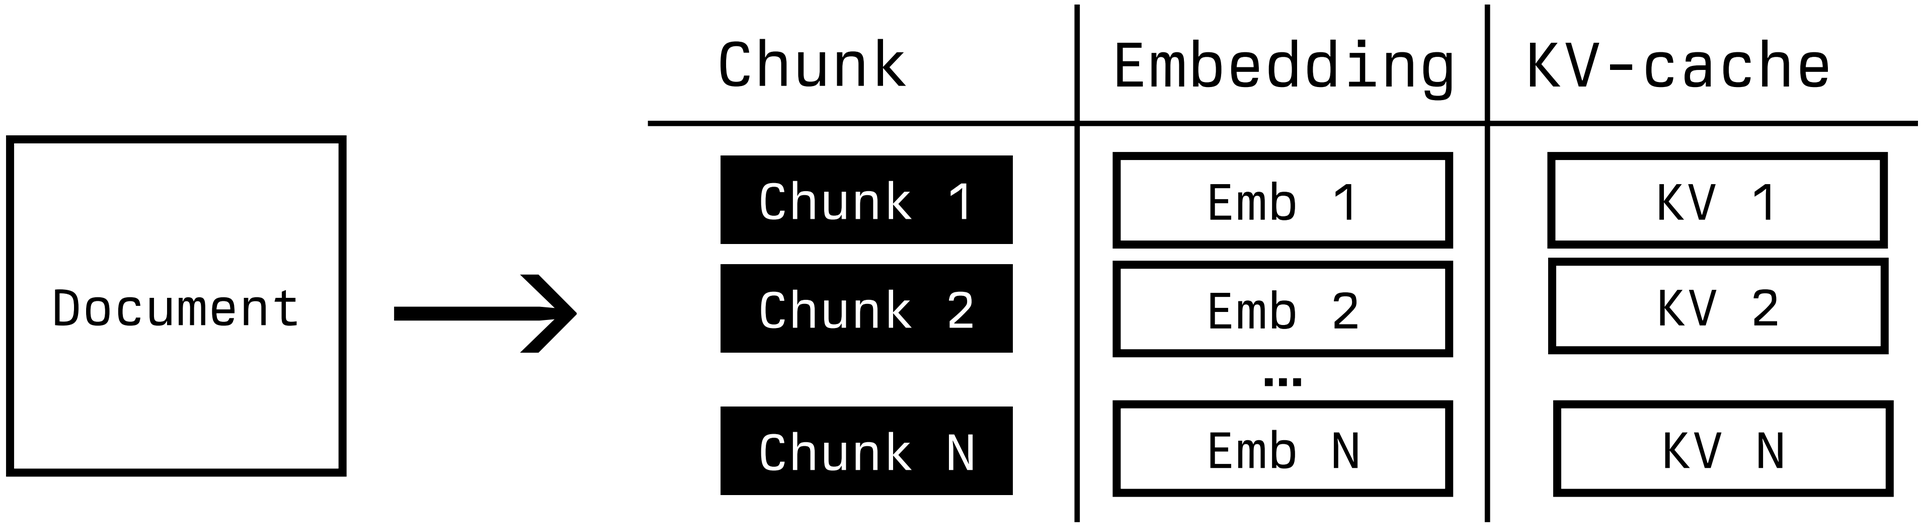

Flash RAG Inference

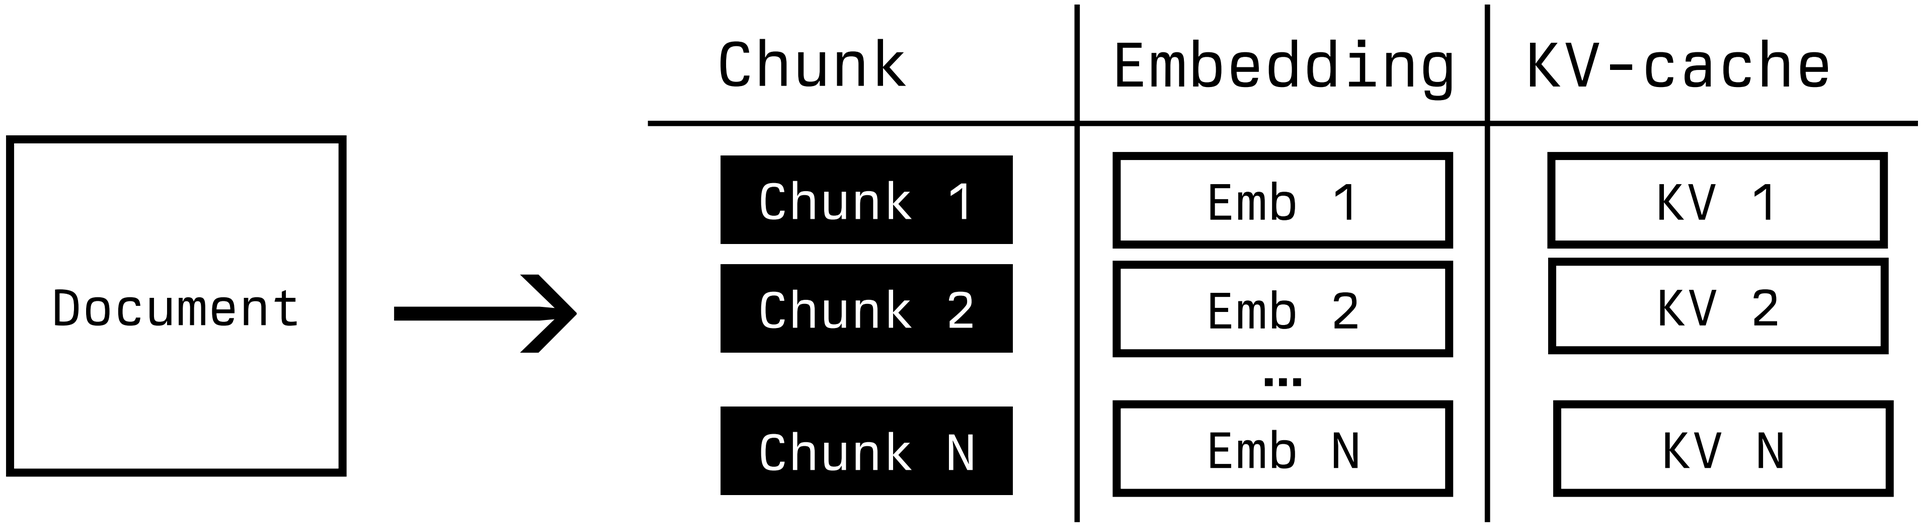

Flash RAG Inference

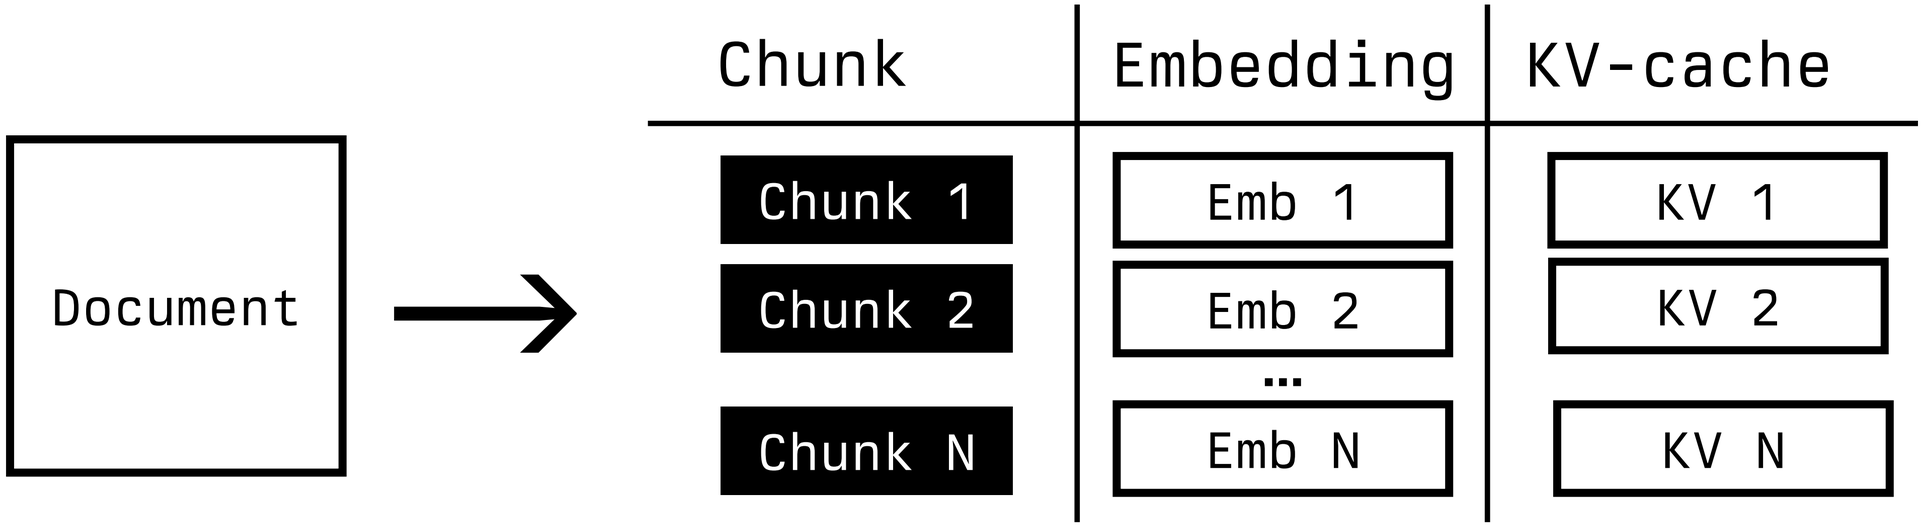

Flash RAG Inference

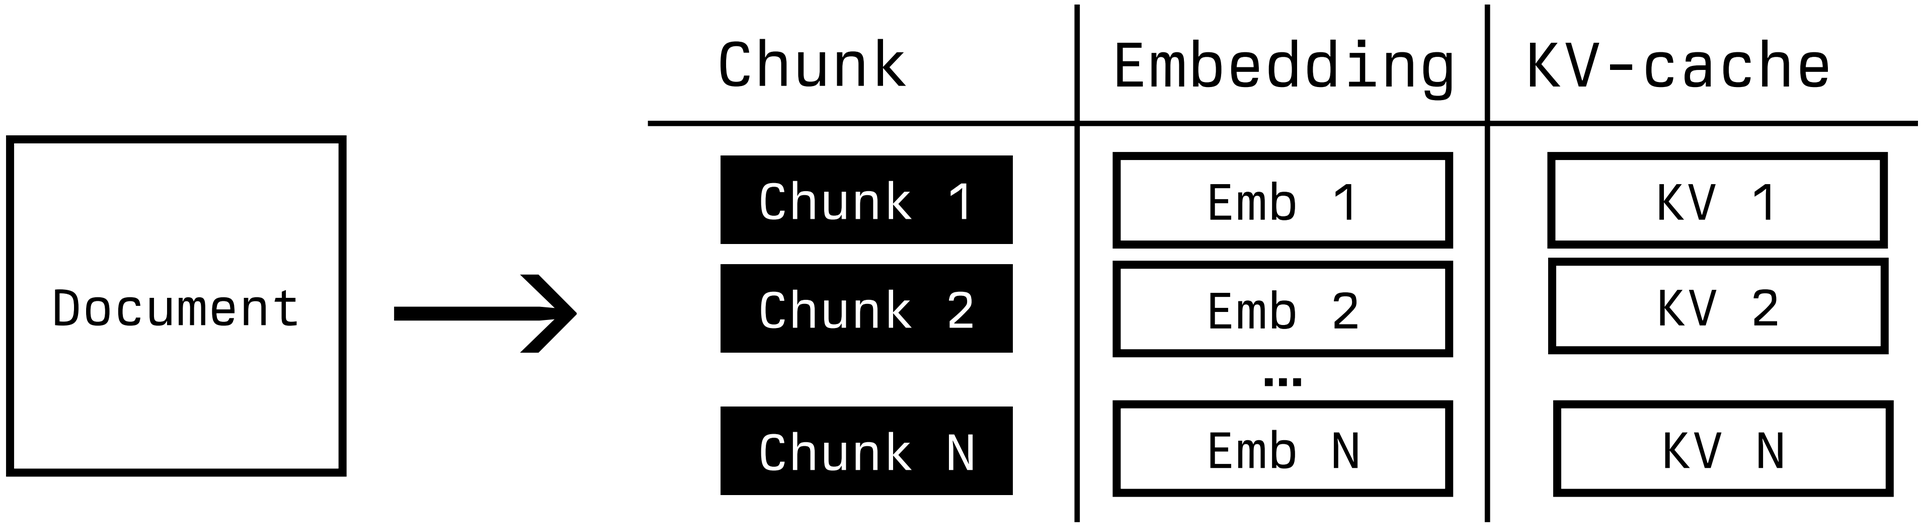

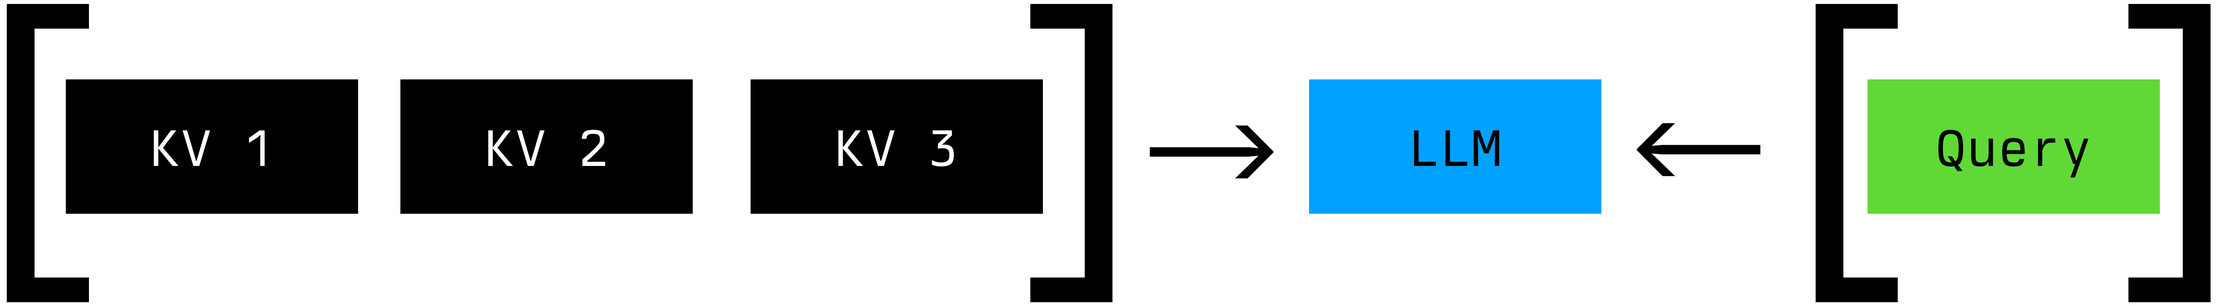

Why is it possible?

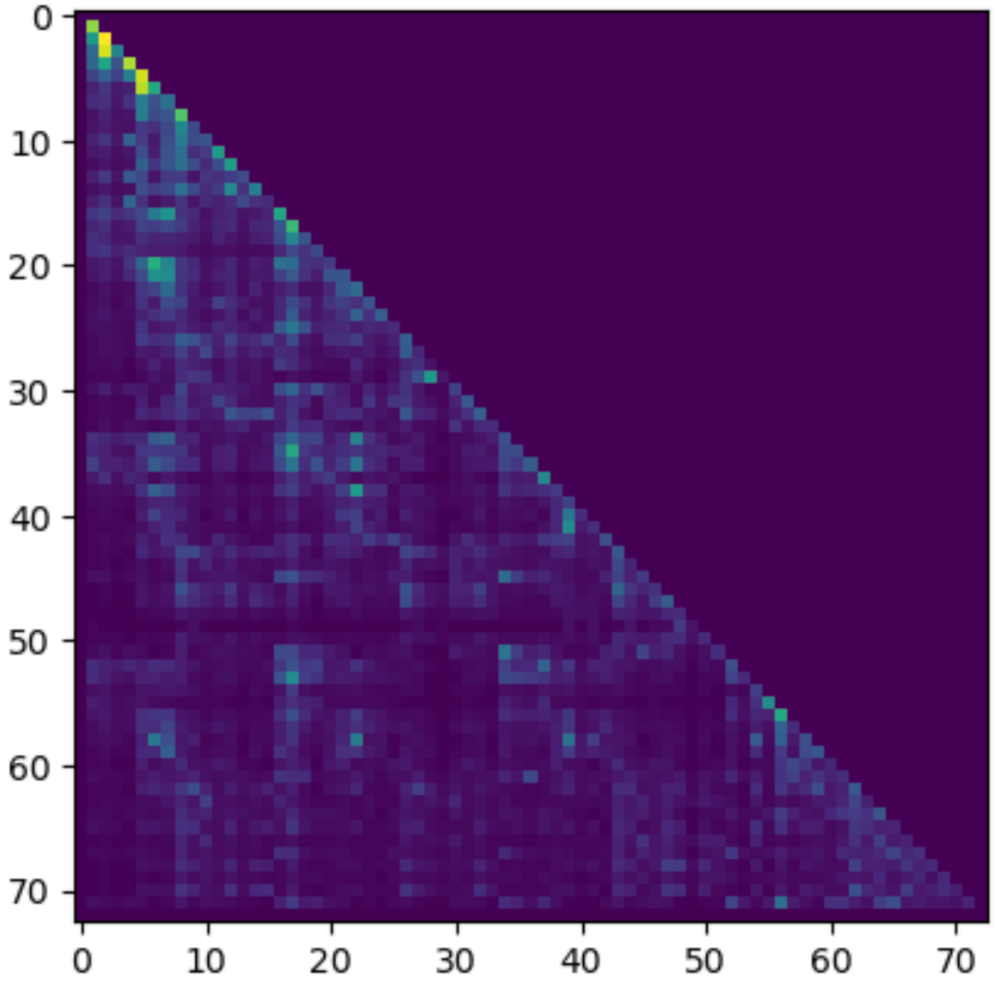

Chunk 1

Chunk 2

Chunk 3

Query

Why is it possible?

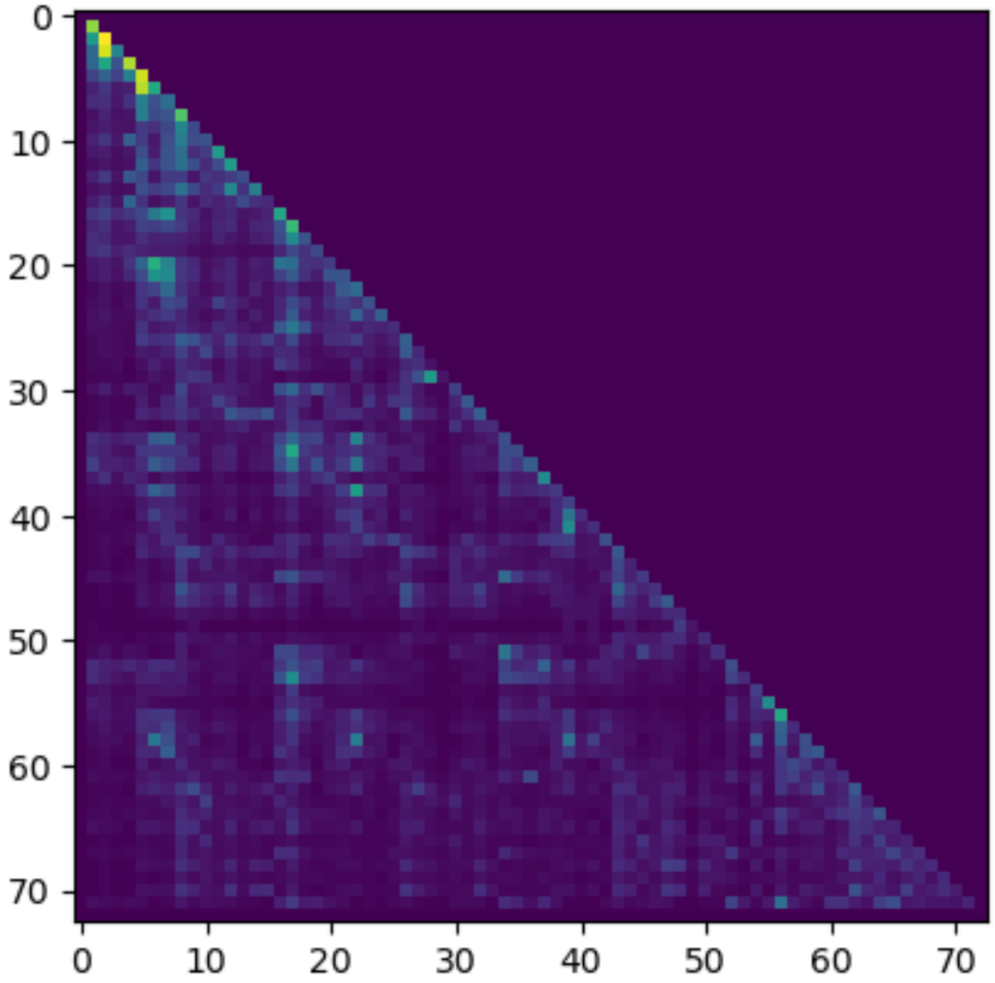

Chunk 1

Chunk 2

Chunk 3

Query

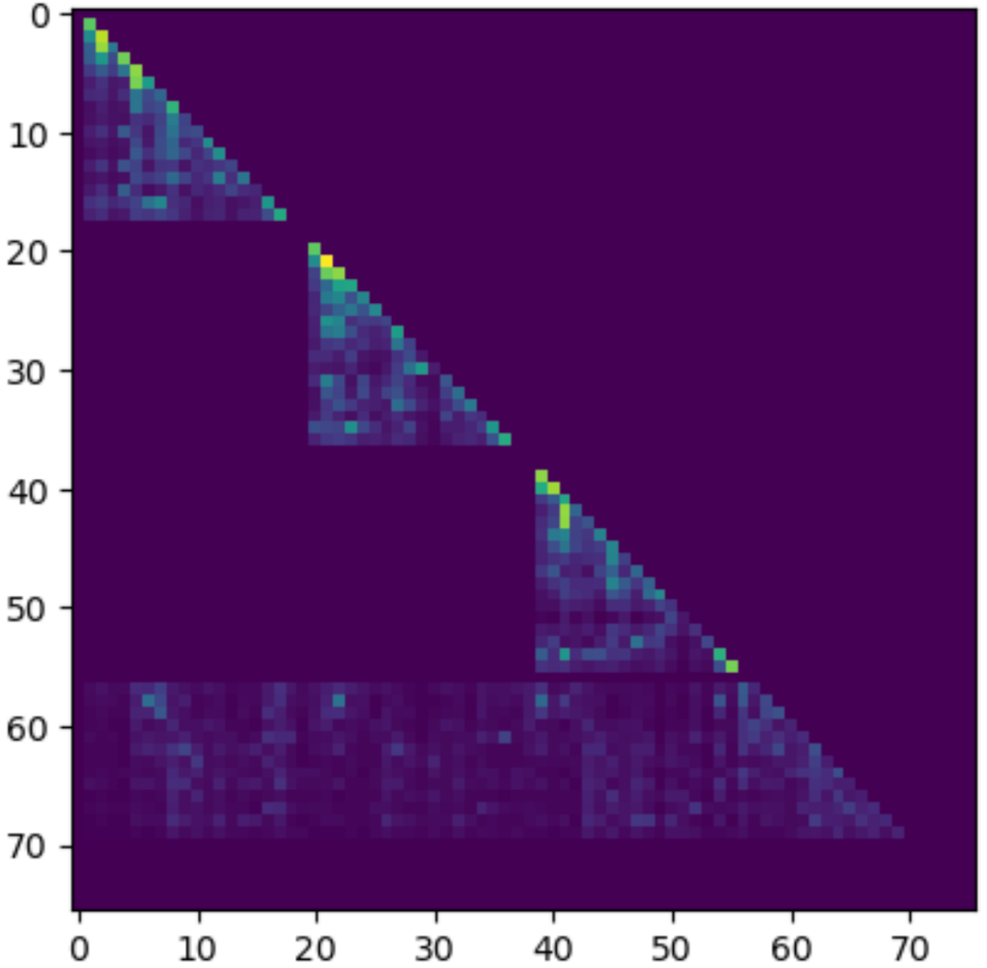

Why is it possible?